

Derrek Gass: 219615449
Alexander Lee: 212490812
Jimmy Le: 217837608
CSC180 - 02 : Haiquan Chen
10-11-2019

Mini-Project 2: Network Intrusion Detector

[Live Demo](#)

(1) Problem Statement

In this project we are looking to implement a predictive model that detects good and bad connections on a computer network to better protect the network from attack. We want to build, train, and optimize our classification model to best generalize the KDD Cup 1999 dataset to be used as a predictive model for future data. In simple terms, we would like to be able to predict if an incoming network connection is an intruder or a normal connection.

(2) Methodology

This problem is a classification type problem, and we utilize several tools available with Python. First we downloaded the KDD cup dataset and reviewed the different columns. The dataset had a large number of redundant records, so we dropped those. We one-hot-coded the categorical column for 'outcome' and created a balanced randomly-selected dataset between good and bad connections. This enables the categorization algorithm to learn to distinguish the two groups equally.

After cleaning the dataset, we arrived at an input array X with 118 parameters, and a y binary output that corresponds to "good" or "bad." In total, we analyzed a balanced dataset of 105,508 records. Splitting this into $\frac{3}{4}$ and $\frac{1}{4}$ created a training set of 86,631 records and 18,877 test records. For Tensorflow to work with a 2-D Convolutional Neural Network, one additional step had to be performed to reshape the X numpy array into the correct number of dimensions for analysis.

(3) Experimental Results and Analysis

We arrived at our best solution with the Convolutional Network over the fully connected Dense network. The CNN took longer than the Dense network to compute weights, but training with a GPU increased the calculation speed.

Our Convolutional model did the best with a Conv2D layer with 64 nodes across, and a kernel size of 1x1, and tanh activation.. This was fed into a MaxPooling2D layer, then into another Conv2D layer the size of the number of X's columns (118) with another tanh activation. This was also maxPooled, then flattened, and fed into a dense network with relu activation, a dropout layer to prevent overfitting, and into the final softmax-activated layer. The optimizer with the best results was adam in this case.

Further analysis of the performance and detailed description of our model is listed in the graphics below.

Best Convolutional Neural Network:

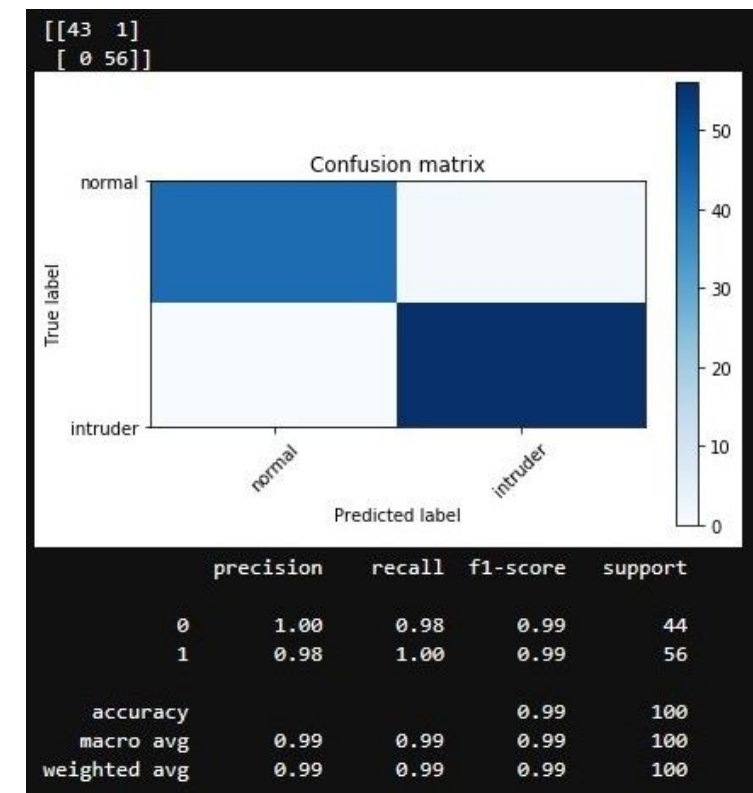
Model: "sequential_20"

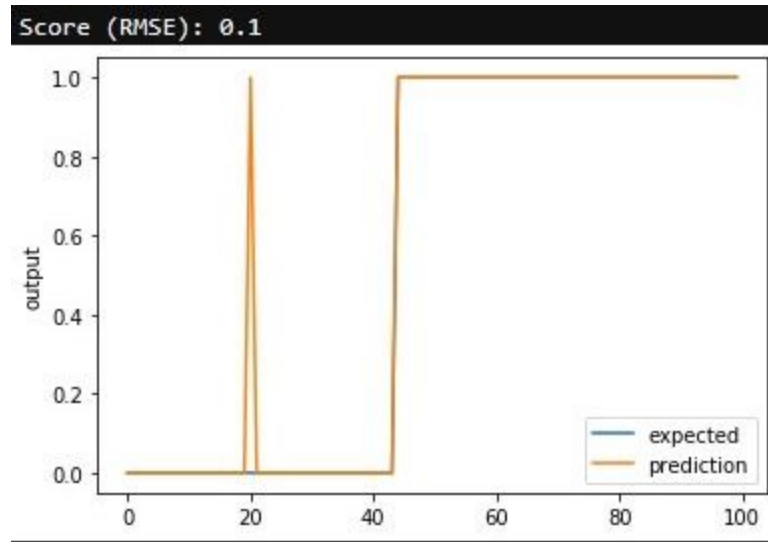
Layer (type)	Output Shape	Param #
conv2d_10 (Conv2D)	(None, 1, 117, 64)	192
max_pooling2d_10 (MaxPooling)	(None, 1, 58, 64)	0
conv2d_11 (Conv2D)	(None, 1, 57, 118)	15222
max_pooling2d_11 (MaxPooling)	(None, 1, 28, 118)	0
flatten_5 (Flatten)	(None, 3304)	0
dense_43 (Dense)	(None, 32)	105760
dropout_5 (Dropout)	(None, 32)	0
dense_44 (Dense)	(None, 2)	66

Total params: 121,240
Trainable params: 121,240
Non-trainable params: 0

86631/86631 - 5s - loss: 0.0076 - acc: 0.9976 - val_loss: 0.0059 - val_acc: 0.9985
Epoch 00012: early stopping
Model: CNN, Activation: tanh, tanh, relu, softmax, Optimizer: adam, Kernel number/size: 118/64, 1x1
Elapsed time: 0:01:04.74

Test loss: 0.021198396283434704
Test accuracy: 0.9900000095367432





Best Fully-Connected Dense Network:

Model: "sequential_22"

Layer (type)	Output Shape	Param #
dense_48 (Dense)	(None, 32)	3808
dense_49 (Dense)	(None, 6)	198
dense_50 (Dense)	(None, 2)	14
Total params: 4,020		
Trainable params: 4,020		
Non-trainable params: 0		

Epoch 10/100

86631/86631 - 6s - loss: 0.0097 - acc: 0.9971 - val_loss: 0.0087 - val_acc: 0.9982

Epoch 00010: early stopping


Model: Sequential, Activation: tanh, Layers, count: 32, 6, 2 Optimizer: adam

Elapsed time: 0:19:30.26

Test loss: 0.050018357187509535


Test accuracy: 0.9900000095367432

```
[[43  1]
 [ 0 56]]
```



	precision	recall	f1-score	support
0	1.00	0.98	0.99	44
1	0.98	1.00	0.99	56
accuracy			0.99	100
macro avg	0.99	0.99	0.99	100
weighted avg	0.99	0.99	0.99	100

Score (RMSE): 0.1



Not-so-good models:

CNN:

```
Epoch 23/100  
86631/86631 - 7s - loss: 0.0153 - acc: 0.9951 - val_loss: 0.0066 - val_acc: 0.9983  
Epoch 00023: early stopping  
Model: CNN, Activation: tanh, sigmoid, relu, softmax, Optimizer: adam, Kernel number/size: 118, 1x1  
Elapsed time: 0:02:51.61
```

```
Epoch 12/100  
86631/86631 - 7s - loss: 0.0078 - acc: 0.9978 - val_loss: 0.0059 - val_acc: 0.9981  
Epoch 00012: early stopping  
Model: CNN, Activation: relu, sigmoid, relu, softmax, Optimizer: adam, Kernel number/size: 32, 1x1  
Elapsed time: 0:01:31.97
```

Dense:

```
86631/86631 - 6s - loss: 0.0103 - acc: 0.9971 - val_loss: 0.0145 - val_acc: 0.9972  
Epoch 00006: early stopping  
Model: Sequential, Activation: relu, Layers, count: 32, 10, 2 Optimizer: adam  
Elapsed time: 0:22:43.86
```

```
86631/86631 - 6s - loss: 0.0105 - acc: 0.9967 - val_loss: 0.0151 - val_acc: 0.9941  
Epoch 00007: early stopping  
Model: Sequential, Activation: tanh, Layers, count: 32, 10, 2 Optimizer: adam  
Elapsed time: 0:24:24.14
```

```
Epoch 5/100  
86631/86631 - 5s - loss: 0.6883 - acc: 0.5056 - val_loss: 0.7099 - val_acc: 0.5044  
Epoch 00005: early stopping  
Model: Sequential, Activation: relu, Layers, count: 64, 24, 2 Optimizer: sgd  
Elapsed time: 0:16:59.06
```

(4) Task Division and Project Reflection

All tasks were divided equally amongst the three of us, as understanding the minutiae of this aspect of computational science is important to furthering research and application in the process of machine learning and data science. Challenges included aspects of data wrangling, such as dealing with redundant data and organizing categorical columns into the one-hot numpy format TensorFlow needs. Tuning hyperparameters, documentation and comparing models also takes time. Our model would perform even better if the dataset contained a larger number and an equal representation of the variety of network attacks. Some attacks only had 2 records, so may be underrepresented in the current model. Should this later be expanded to detect the type of attack, one would add a layer and create a new y output array with the network attacks one-hot-coded into unique columns and the final layer would have 23 neurons; 22 for the attacks and 1 for normal connections.